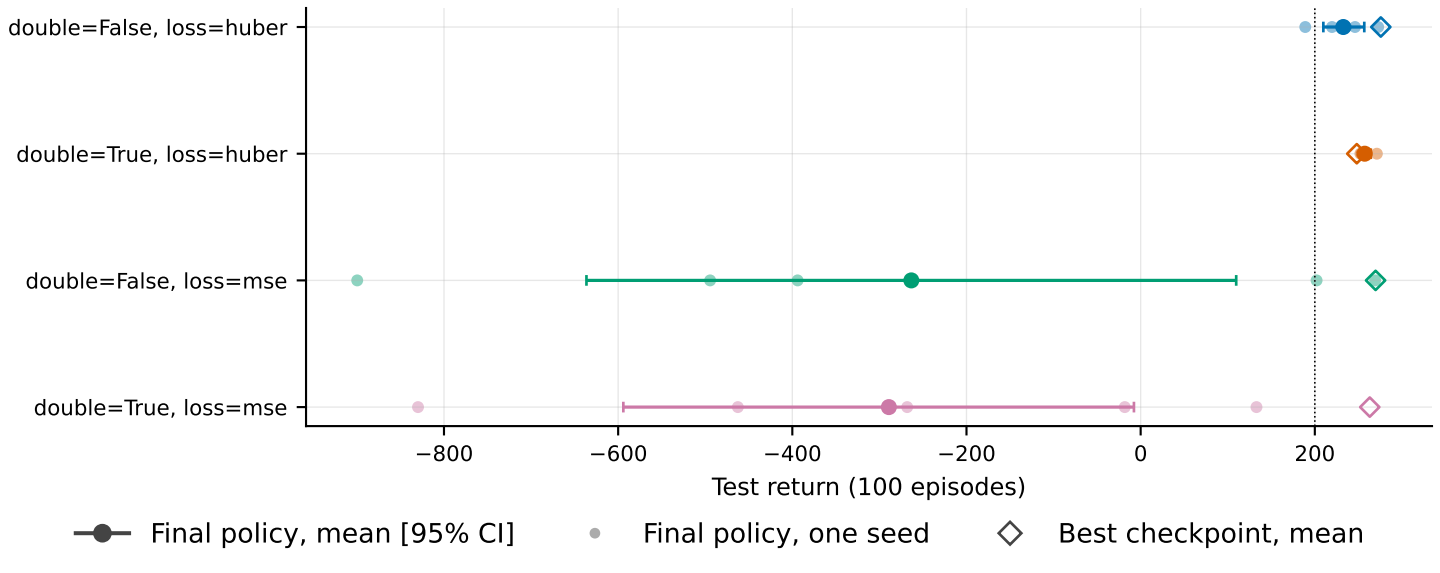
double=False, loss=huber
double=True, loss=huber
double=False, loss=mse
double=True, loss=mse
−800
−600
−400
−200
0
200
Test return (100 episodes)
Final policy, mean [95% CI]
Final policy, one seed
Best checkpoint, mean